

R Lab 0 Exercises

Basic Data Manipulation & Plotting in R

Datasaurus

- A. Load in `datasaurus.csv` (on the Canvas page for R Lab 0; more information can be found at github.com/stephlocke/datasauRus) using the following command (change the file argument as necessary to point to the correct file path for the computer you're using):

```
dino <- read.csv(file = "./Rlab0/datasaurus.csv", header = T)
```

Notes: An alternative function to read in a comma separated file is `read.table()`, adding a `sep = ","` argument. Also, if you have RStudio, note that “dino” now appears within the Environment tab on the upper-right of the screen.

- B. Run the command `head(dino)`, followed by `names(dino)`.
- C. Run the commands `hist(dino$x)`, `boxplot(dino$x)`, and `plot(dino$x)`.
- D. Run the command `plot(dino$x, dino$y)`.
- E. Run the command `?plot`. Use the results to add a title of “T-Rex” and a subtitle of “technology is truly amazing” to your plot in D.
- F. Run the following code, which replaces the first two entries in the `dino$x` vector. Explain in language understandable to an R beginner why the output of the first line differs from the third line.

```
typeof(dino$x)
dino$x[1:2] <- c("hello", "there")
typeof(dino$x)
```

- G. Why does the following command return TRUE?

```
identical(dino$y, dino[, 2])
```

- H. What `typeof()` R vector is returned by the following command?

```
dino$y < 20
```

- I. With the last question in mind, run the following two commands. What do you conclude about ways you can index things in R?

```
dino_subset_1 <- dino[dino$y < 20, ]  
dino_subset_2 <- dino[c(1, 3, 6:10), ]
```

- J. Combining code from Part B and the second line of Part F, change the column names of dino to your first name and last name. Run head(dino) again to make sure that the changes worked.

Traffic Accidents

The file denvertraffic.txt contains traffic accident data from the City of Denver open data portal (denvergov.org/opendata/dataset/city-and-county-of-denver-traffic-accidents) from July 2017.

- A. Read in denvertraffic.txt. Name the resulting data frame traffic.
- B. Print the column names of the dataset. (Revisit the command you used in Part A if you get V1, V2, ..., V20)
- C. Run the following command:

```
dia_only <- traffic[traffic$NEIGHBORHOOD_ID == "DIA", ]  
head(dia_only)
```
- D. Run table(traffic\$OFFENSE_TYPE_ID)
- E. What does the following code do?

```
traffic[traffic$OFFENSE_TYPE_ID == "TRAF-ACCIDENT-",  
       "OFFENSE_TYPE_ID"] <- "TRF-ACCIDENT"
```
- F. Use code in the style of Part C to make a data frame consisting of only the offenses marked as a DUI. Save this data frame as dui_only.
- G. Use plot() to make a scatter plot of all traffic accidents, with longitude on the x-axis and latitude on the y-axis. Add the argument col = "grey" to make the points grey.
- H. If we were to run another plotting command, by default R would overwrite our last plot. To avoid this, try out the command points() instead. points() takes very

similar arguments to `plot()`.

Use `points()` to add the geographical locations of accidents in `dia_only` data with `col = "blue"`. Similarly, add accidents associated with `dui_only` in red.

- I. A pound sign (`#`) can be used in R to specify comments. Run the following commands, then annotate it by writing what each line does.

```
# write comment here
bike_accidents <- table(traffic$NEIGHBORHOOD_ID,
                        traffic$BICYCLE_IND)

# write comment here
length(bike_accidents)

# write comment here
prop <- bike_accidents[,2] / (bike_accidents[,1]+bike_accidents[,2])

# write comment here
max(prop)

# write comment here
names(prop)[ prop == max(prop) ]
```

Note: while these examples include comments on their own line, you can also add comments at the end of a line of code. However, any functions included after the pound sign will not run. For example,

```
max(prop) #comment after the max() function
```

```
max(prop) #in this case the min() function won't run: min(prop)
```